Unified Simplified Grapheme Acoustic Modeling for Medieval Latin LVCSR



Lili Szabó, Péter Mihajlik, András Balog, Tibor Fegyó

What is the problem with Latin speech recognition?

- Latin is not spoken natively
- There is no available speech database, and it is resource-heavy to create one
- Many variants/dialects exists, and we can only make guesses about the pronunciation
- The pronunciation mainly depends on
 - the era of the read text
 - the native language of the speaker

Text data

Regions of origin: Kingdom of Bohemia (CZ), Kingdom of Hungary (HU), Kingdom of Poland (PL)

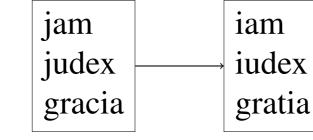
- In-domain data (Monasterium): medieval charters (HU), 480k/35k token/type
- Background data (Latin Library): historical texts, 1.3M/115k token/type

Speech data

Languages: CZ, HU, PL, RO

gratia

Spelling variants



Test data

- Independent medieval charters
- Region of read text: CZ, HU, PL
- Native language of test speakers: CZ, HU, PL, SK

Perplexity measures on test

Table 1: Perplexity/OOV rate

	Te	Text region				
Corpus	CZ	HU	PL	All		
Monasterium	551	82	3130	671		
Latin Library	3266	3549	2305	4303		
Interpolated	924	82	2288	953		

Acoustic model

Language model

- 3-gram language model
- Kneser-Ney smoothing
- Interpolating the two corpora
- SRILM [2]

- Mel-Frequency Cepstrum + Energy features were used with Linear Discriminant Analysis (LDA) + Maximum Likelihood Linear Transformation (MLLT), with a splice context of ± 4 frames, 10 ms of frame shift.
- \bullet 9×40 dimensional spliced up feature vectors served as input to the feed-forward, 6 hidden-layer neural network with p-norm [1] activation function.
- Prior to DNN training, a Gaussian Mixture Model (GMM) pre-training was performed.
- Clustering and Regression Tree (CART) [1] was applied to obtain acrossword context dependent shared state phone (or graph) models and their time alignment.
- The number of senones (and so the size of the DNN softmax output layer) was between 7.000 and 11.000 depending on the nature of the training data.
- The size of the hidden layers was kept constantly on 2.000.
- A minibatch size of 512, an initial learning rate of 0.1, and final learning rate of 0.01 was applied in 20 epochs using the Kaldi toolkit [1].

Baseline Grapheme Model

Languages: Czech (CZ), Hungarian (HU), Polish (PL), Romanian (RO)

- All graphemes are trained
- Only those grapheme models are retained that are part of the Latin alphabet

Table 2: Word Error Rate (WER[%]) results for monolingual grapheme-based acoustic models of Czech, Hungarian, Polish and Romanian (CZ, HU, PL, RO).

	S	Speaker				
AM Language	CZ	HU	PL	SK	\sum	
CZ	53.6	73.8	62.9	45.7	59.0	
HU	33.7	28.6	47.1	29.1	34.6	
PL	65.0	67.6	46.4	51.1	57.5	
RO	53.6	69.1	44.7	43.8	52.8	

Knowledge-based grapheme-to-phoneme (G2P) mapping

Languages: CZ, HU

Table 3: Latin digraph context-insensitive rewrite rules.

		Digraph				
		ae	oe	ph	qu	
C'	Z	e	oe	f	kv	
H	U	e	Ø	f	kv	

Table 4: Latin context-sensitive rewrite rules. V: vowel, VP: palatal vowel, ^VP: everything but a palatal vowel, C: consonant, *: zero or any, $\hat{}$: beginning of word, $\hat{}$ and $\hat{}$: not s, t or x.

GR	c	С	ch	ch	gu	gu	ti	ti
					_	_	tsi	
rule	cVP	c^VP	VC*ch	^C*ch	guV	guC	$[\hat{s}tx]$ tiV	tiC

Table 5: WER[%] for Czech-Latin sourcetarget G2P model. Acoustic model training set: 76 hours.

	Latin Test Text					
Speaker	CZ	HU	PL	\sum		
CZ	43.8	28.2	49.1 58.7	40.4		
HU						
PL	53.3	18.2	53.2	41.6		
SK	30.3	30.0	44.0	34.8		
$\overline{\sum}$	43.9	28.9	50.8	41.2		

Table 6: WER[%] for Hungarian-Latin source-target G2P model. Acoustic model training set: 567 hours.

	Latir	n Test	Text	
Speaker	CZ	HU	PL	\sum
CZ	19.4	6.4	28.0	17.9
HU	25.0	25.4	20.2	23.5
PL	28.9	15.4	41.3	28.5
SK	20.4	9.1	22.9	17.5
\sum	22.6	12.5	28.1	21.1

Unified Simplified Grapheme (USG) Model

Languages: CZ, HU, PL, RO

Table 7: Simplification examples for the unified model.

Language	CZ	HU	PL	RC
Orthographic form	řekl	őz	miś	apa
USG transcription	rekl	ΟZ	mis	apa

Table 8: WER[%] for all the three-language USG models.

Speaker	
AM Language CZ HU PL SK	\sum
CZ+HU+PL 28.2 28.2 27.7 22.4 2	26.6
CZ+HU+RO 23.3 21.4 23.9 19.2 2	21.9
CZ+PL+RO 24.6 33.1 25.6 19.8 2	25.8
HU+PL+RO 24.8 21.5 25.7 20.7 2	23.2

WER[%] for USG model of Table 9: Czech, Hungarian, Polish and Romanian (CZ+HU+PL+RO).

	Latır	1 Test	Text	
Speaker	CZ	HU	PL	\sum
CZ		11.8		
HU	21.1	14.6	25.7	20.5
PL	23.0	10.0	33.0	22.0
SK	14.5	12.7	24.8	17.3
\sum	19.9	12.2	29.0	20.4

Speech data: CZ, HU, PL, RO Model type: baseline, knowledge-based, USG

Speaker

Dimensions of data

Region of read text: CZ, HU, PL

Native language of test speakers: CZ, HU, PL, SK

System diagram USG Text Base Acoustic Language Model Model HU Medieval Latin ASR SK Evaluate

Read text

Conclusions

- Four-language USG is the best
- It is able to generalize over different speaker test sets

References

- [1] Povey, D., Ghoshal, A., Boulianne, G., Burget, L., Glembek, O., Goel, N., Hannemann, M., Motlicek, P., Qian, Y., Schwarz, P., Silovsky, J., Stemmer, G., Vesely, K.: The kaldi speech recognition toolkit. In: IEEE 2011 Workshop on Automatic Speech Recognition and Understanding. IEEE Signal Processing Society (2011)
- [2] Stolcke, A.: Srilm an extensible language modeling toolkit. In: In Proceedings of the 7th International Conference on Spoken Language Processing (ICSLP). pp. 901–904 (2002)

Figure 1: Medieval Latin Speech Recognizer